

# Practical R: Data Ingestion and Munging

Abhijit Dasgupta

Fall, 2019

# A quick refresh

- We talked about various data structures in R
- The primacy of the `data.frame`
  - Extracting individual variables from a data frame
    - `breast_cancer$ER.Status`, `breast_cancer[, 'ER.Status']`,  
`breast_cancer[['ER.Status']]`
  - Extracting rows of a `data.frame`
- Identifying data classes using the `class` function
- Recognizing different classes: `numeric`, `character`, `factor`, `Date`, ..
  - testing for a class: `is.numeric`
  - converting to a class: `as.numeric`

# **A note on factors**

# Factors

- Factors are stored internally as integers, with *meta-data* in the form of text labels
  - There is an inherent ordering of labels, by default alphabetically
- Individual levels of a factor are treated as *separate* but related variables (dummy variables)

```
breast_cancer <- read_csv('data/clinical_data_breast_cancer_modified.csv')
names(breast_cancer) <- make.names(names(breast_cancer))
breast_cancer$ER.Status.f <- factor(breast_cancer$ER.Status)
summary(breast_cancer$ER.Status)
```

```
#>      Length      Class      Mode
#>      105 character character
```

```
summary(breast_cancer$ER.Status.f)
```

```
#> Indeterminate      Negative      Positive
#>              1             36             68
```

# Factors

```
breast_cancer$ER.Status.f <- fct_relevel(breast_cancer$ER.Status.f, 'Negative')  
summary(breast_cancer$ER.Status.f)
```

```
#>      Negative Indeterminate      Positive  
#>          36             1          68
```

This is manipulating the meta-data, not the actual data itself

# Factors

```
breast_cancer$ER.Status.n <- as.numeric(breast_cancer$ER.Status.f)
summary(breast_cancer$ER.Status.n)
```

```
#>      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
#>      1.000   1.000   3.000   2.305   3.000   3.000
```

Logistic regression of death status on ER status

```
#> # A tibble: 2 x 2
#>   term      estimate
#>   <chr>      <dbl>
#> 1 (Intercept)    1.81
#> 2 ER.Status.n     0.148
```

Only one coefficient, since levels are modeled as numeric, with one slope being estimated

```
#> # A tibble: 3 x 2
#>   term      estimate
#>   <chr>      <dbl>
#> 1 (Intercept)    2.08
#> 2 ER.Status.fIndeterminate -17.6
#> 3 ER.Status.fPositive      0.256
```

One coefficient for all but one factor level

# RMarkdown tip of the day

You can add options to each R chunk to add or suppress output

Option	Property
echo=T/F	Does the document show the R code
eval=T/F	Does the chunk get evaluated by R
message=T/F	Do messages get printed
warning=T/F	Do warnings get printed

You can also set these once per session by putting the following in a R chunk:

```
knitr::opts_chunk(echo=T, eval=T, message=F, warning=F)
```

See [here](#) for the full gory details

# Data ingestion



# Data ingestion

Unlike Excel, you have to pull data into R for R to operate on it

Typically your data is in some sort of file (Excel, csv, sas7bdat, dta, txt)

You need to find a way to pull it into R

The GUI you've used is one way, but not very programmatic

# Data ingestion

Type	Function	Package	Notes
csv	read_csv	readr	Takes care of formatting
csv	read.csv	base	Built in
csv	fread	data.table	Fastest
Excel	read_excel	readxl	
sas7bdat	read_sas	haven	SAS format
sav	read_spss	haven	SPSS format
dta	read_dta	haven	Stata format

# Data ingestion

We will use [this](#) csv data and [this](#) Excel data for the following:

```
brca_clinical <- readr::read_csv('data/BreastCancer_Clinical.csv')
brca_clinical2 <- data.table::fread('data/BreastCancer_Clinical.csv')
```

```
str(brca_clinical)
```

```
#> Classes 'spec_tbl_df', 'tbl_df', 'tbl' and 'data.frame': 105 obs
#> $ Complete TCGA ID      : chr "TCGA-06-0001-01A-1101-1001"
#> $ Gender                 : chr "FEM"
#> $ Age at Initial Pathologic Diagnosis: num 66.4
#> $ ER Status              : chr "Neg"
#> $ PR Status              : chr "Neg"
#> $ HER2 Final Status      : chr "Neg"
#> $ Tumor                  : chr "T3"
#> $ Tumor--T1 Coded        : chr "T_0"
#> $ Node                   : chr "N3"
#> $ Node-Coded             : chr "Pos"
#> $ Metastasis             : chr "M1"
#> $ Metastasis-Coded       : chr "Pos"
#> $ AJCC Stage             : chr "Sta"
#> $ Converted Stage        : chr "No"
#> $ Survival Data Form     : chr "fol"
#> $ Vital Status           : chr "DEC"
#> $ Days to Date of Last Contact : num 240
```

```
str(brca_clinical2)
```

```
#> Classes 'data.table' and 'data.frame': 105 obs
#> $ Complete TCGA ID      : chr "TCGA-06-0001-01A-1101-1001"
#> $ Gender                 : chr "FEM"
#> $ Age at Initial Pathologic Diagnosis: int 66.4
#> $ ER Status              : chr "Neg"
#> $ PR Status              : chr "Neg"
#> $ HER2 Final Status      : chr "Neg"
#> $ Tumor                  : chr "T3"
#> $ Tumor--T1 Coded        : chr "T_0"
#> $ Node                   : chr "N3"
#> $ Node-Coded             : chr "Pos"
#> $ Metastasis             : chr "M1"
#> $ Metastasis-Coded       : chr "Pos"
#> $ AJCC Stage             : chr "Sta"
#> $ Converted Stage        : chr "No"
#> $ Survival Data Form     : chr "fol"
#> $ Vital Status           : chr "DEC"
#> $ Days to Date of Last Contact : int 240
```

# A note on two "super"-data.frame objects

## A tibble

```
#> # A tibble: 6 x 30
#>   `Complete TCGA ...` Gender `Age at Initial...` `ER St
#>   <chr>                <chr>                <dbl> <chr>
#> 1 TCGA-A2-A0T2        FEMALE                66 Negati
#> 2 TCGA-A2-A0CM        FEMALE                40 Negati
#> 3 TCGA-BH-A18V        FEMALE                48 Negati
#> 4 TCGA-BH-A18Q        FEMALE                56 Negati
#> 5 TCGA-BH-A0E0        FEMALE                38 Negati
#> 6 TCGA-A7-A0CE        FEMALE                57 Negati
#> # ... with 25 more variables: `HER2 Final Status` <chr>,
#> #   `Tumor--T1 Coded` <chr>, `Node` <chr>, `Node-Coded` <chr>,
#> #   `Metastasis` <chr>, `Metastasis-Coded` <chr>, `
#> #   `Converted Stage` <chr>, `Survival Data Form` <chr>,
#> #   `Status` <chr>, `Days to Date of Last Contact` <chr>,
#> #   `Death` <dbl>, `OS event` <dbl>, `OS Time` <dbl>,
#> #   `SigClust Unsupervised mRNA` <dbl>, `SigClust
#> #   `miRNA Clusters` <dbl>, `methylation Clusters
#> #   `Clusters` <chr>, `CN Clusters` <dbl>, `Integr
#> #   `PAM50` <dbl>, `Integrated Clusters (no exp)` <dbl>,
#> #   `Integrated Clusters (unsup exp)` <dbl>
```

## A data.table

```
#>   Complete TCGA ID Gender Age at Initial Patholo
#> 1:   TCGA-A2-A0T2 FEMALE
#> 2:   TCGA-A2-A0CM FEMALE
#> 3:   TCGA-BH-A18V FEMALE
#> 4:   TCGA-BH-A18Q FEMALE
#> 5:   TCGA-BH-A0E0 FEMALE
#> 6:   TCGA-A7-A0CE FEMALE
#>   PR Status HER2 Final Status Tumor Tumor--T1 Co
#> 1: Negative                Negative T3      T_Ot
#> 2: Negative                Negative T2      T_Ot
#> 3: Negative                Negative T2      T_Ot
#> 4: Negative                Negative T2      T_Ot
#> 5: Negative                Negative T3      T_Ot
#> 6: Negative                Negative T2      T_Ot
#>   Metastasis Metastasis-Coded AJCC Stage Convert
#> 1: M1          Positive      Stage IV   No_Co
#> 2: M0          Negative      Stage IIA   S
#> 3: M0          Negative      Stage IIB   No_Co
#> 4: M0          Negative      Stage IIB   No_Co
#> 5: M0          Negative      Stage IIIC  No_Co
#> 6: M0          Negative      Stage IIA   S
#>   Survival Data Form Vital Status Days to Date o
#> 1:      followup      DECEASED
#> 2:      followup      DECEASED
#> 3:      enrollment      DECEASED
```

# A note on two "super"-data.frame objects

- A `tibble` works pretty much like any `data.frame`, but the printing is a little saner
- A `data.table` is faster, has more inherent functionality, but has a ver different syntax

We'll work almost entirely with `tibble`'s and not `data.table`

Suggested modifications:

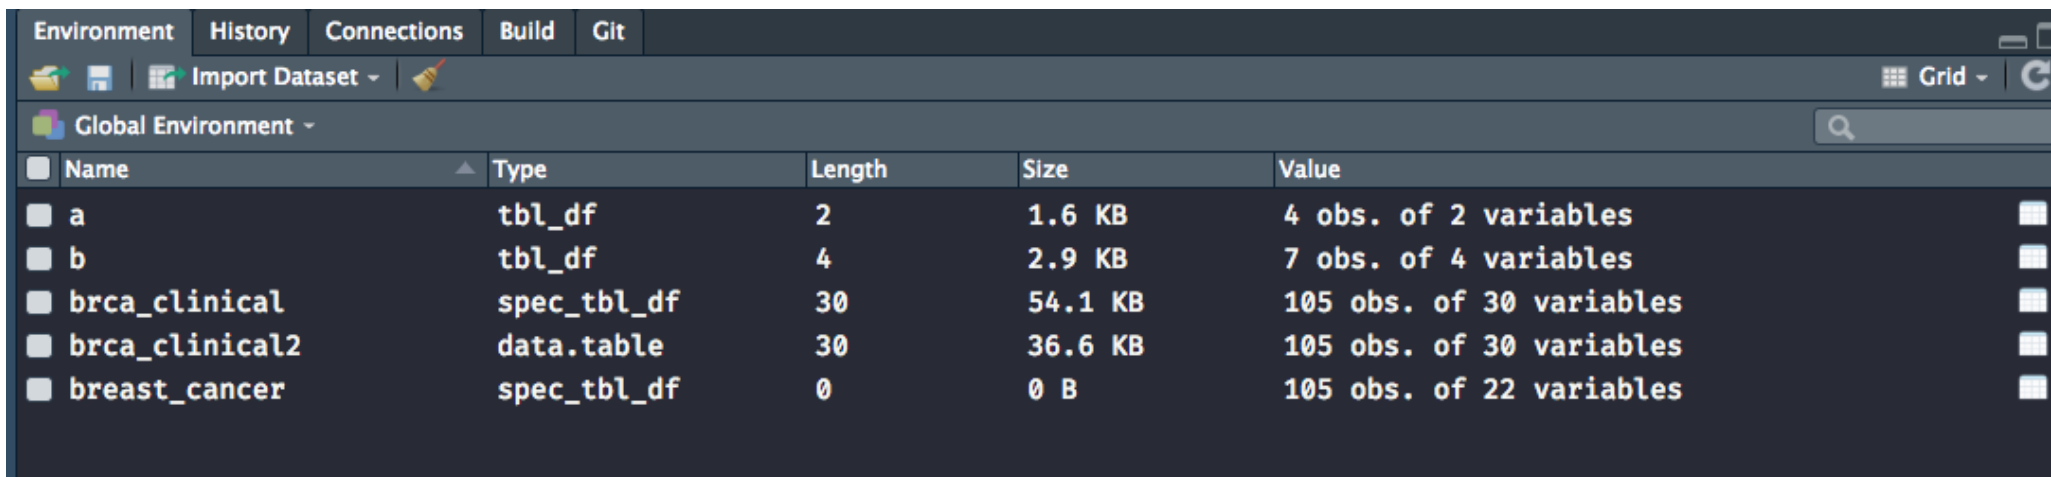
- If using `fread`, convert the resulting object to a `data.frame` or `tibble` using `as_data_frame()` or `as_tibble`
- Convert the column names to not have spaces using, for example,

```
names(brca_clinical) <- make.names(names(brca_clinical))
```

# Data ingestion

Note that you **have** to give a name to what you're importing using `read_*` or whatever you're using, otherwise it won't stay in R

```
brca_clinical <- readr::read_csv('data/BreastCancer_Clinical.csv')
```



The screenshot shows the RStudio Environment pane. At the top are tabs for Environment, History, Connections, Build, and Git. Below the tabs is a toolbar with icons for file operations and a search bar. The main area displays the 'Global Environment' with a table of objects. The table has columns for Name, Type, Length, Size, and Value. The objects listed are 'a', 'b', 'brca\_clinical', 'brca\_clinical2', and 'breast\_cancer'.

Name	Type	Length	Size	Value
a	tbl_df	2	1.6 KB	4 obs. of 2 variables
b	tbl_df	4	2.9 KB	7 obs. of 4 variables
brca_clinical	spec_tbl_df	30	54.1 KB	105 obs. of 30 variables
brca_clinical2	data.table	30	36.6 KB	105 obs. of 30 variables
breast_cancer	spec_tbl_df	0	0 B	105 obs. of 22 variables

# Reading Excel

```
excel_sheets('data/BreastCancer.xlsx')
```

```
#> [1] "Clinical" "Expression"
```

```
brca_expression <- readxl::read_excel('data/BreastCancer.xlsx', sheet='Expression')
```

# Data export



# Data export

Type	Function	Package	Notes
csv	write_csv	readr	Takes care of formatting
csv	write.csv	base	Built in
csv	fwrite	data.table	Fastest
Excel	write.xlsx	openxlsx	
sas7bdat	write_sas	haven	SAS format
sav	write_spss	haven	SPSS format
dta	write_dta	haven	Stata format

We'll often save tabular results using these functions

# Simplifying import/export

We'll be using a package that makes this easier.

It's called `rio` and it has two basic functions: `import` and `export`.

The `rio` package uses the different packages mentioned earlier but unifies it into a single syntax

## Classwork

Open an Rmd file, and create a R chunk where you use the function `import` from `rio` to load the clinical breast cancer data into R

- Note, you have to "activate" the `rio` package in the chunk
- You have to save the imported object by giving it a name

10:00

# Data munging



# The tidyverse

# What is the tidyverse?

The tidyverse is an opinionated collection of R packages designed for data science. All packages share an underlying design philosophy, grammar, and data structures. -- Dr. Hadley Wickham

- A human-friendly syntax and semantics to make code more understandable
- The functions in the tidyverse often wraps harder-to-understand functions into simpler, more understandable forms
- We're taking an opinionated choice here
  - Covers maybe 85% of the cases you'll ever face
  - Takes a particular viewpoint about how data *should* be organized
- But this makes things easier and simpler

# What's tidy here?

The way data is organized in a data frame is **tidy** in this framework.

1. Each variable must have its own column.
2. Each observation must have its own row.
3. Each value must have its own cell.

In practical terms:

1. Put data in a data frame / *tibble*
2. Make sure each variable is in its own column

# Tidy data

A first step in the tidyverse is to activate the tidyverse meta-package

```
library(tidyverse)
```

- **ggplot2**: Create Elegant Data Visualisations Using the Grammar of Graphics
- **purrr**: Functional Programming Tools
- **readr**: Read Rectangular Text Data
- **tidyr**: Easily Tidy Data with 'spread()' and 'gather()' Functions
- **dplyr**: A Grammar of Data Manipulation
- **forcats**: Tools for Working with Categorical Variables (Factors)
- **lubridate**: Make Dealing with Dates a Little Easier
- **stringr**: Simple, Consistent Wrappers for Common String Operations



# Tidy data

The common feature of all these packages is that their functions take a data frame (which the tidyverse calls a `tibble`) as their first argument.

So the starting point for any analysis is the data set.

# Tidy data

```
table1
```

```
#> # A tibble: 6 x 4
#>   country    year cases population
#>   <chr>    <int> <int>      <int>
#> 1 Afghanistan 1999    745  19987071
#> 2 Afghanistan 2000   2666  20595360
#> 3 Brazil      1999  37737  172006362
#> 4 Brazil      2000  80488  174504898
#> 5 China       1999 212258 1272915272
#> 6 China       2000 213766 1280428583
```

Is this tidy?

# Tidy data

```
table2
```

```
#> # A tibble: 12 x 4
#>   country      year type      count
#>   <chr>      <int> <chr>      <int>
#> 1 Afghanistan 1999 cases         745
#> 2 Afghanistan 1999 population 19987071
#> 3 Afghanistan 2000 cases         2666
#> 4 Afghanistan 2000 population 20595360
#> 5 Brazil      1999 cases         37737
#> 6 Brazil      1999 population 172006362
#> 7 Brazil      2000 cases         80488
#> 8 Brazil      2000 population 174504898
#> 9 China       1999 cases         212258
#> 10 China      1999 population 1272915272
#> 11 China      2000 cases         213766
#> 12 China      2000 population 1280428583
```

Is this tidy?

# Tidy data

```
table3
```

```
#> # A tibble: 6 x 3
#>   country      year rate
#> * <chr>      <int> <chr>
#> 1 Afghanistan 1999 745/19987071
#> 2 Afghanistan 2000 2666/20595360
#> 3 Brazil       1999 37737/172006362
#> 4 Brazil       2000 80488/174504898
#> 5 China        1999 212258/1272915272
#> 6 China        2000 213766/1280428583
```

Is this tidy?

# Tidy data

```
table4a # cases
```

```
#> # A tibble: 3 x 3
#>   country    `1999` `2000`
#> * <chr>      <int> <int>
#> 1 Afghanistan     745   2666
#> 2 Brazil        37737  80488
#> 3 China         212258 213766
```

```
table4b # population
```

```
#> # A tibble: 3 x 3
#>   country      `1999`      `2000`
#> * <chr>        <int>      <int>
#> 1 Afghanistan 19987071  20595360
#> 2 Brazil      172006362 174504898
#> 3 China       1272915272 1280428583
```

Are these tidy?

# Can we make datasets tidy?

Sometimes. The functions in the `tidyr` package can help

- `separate` is a function that can split a column into multiple columns
  - When there are multiple variables together in a column

```
table3
```

```
#> # A tibble: 6 x 3  
#>   country    year rate  
#> * <chr>    <int> <chr>  
#> 1 Afghanistan 1999 745/19987071  
#> 2 Afghanistan 2000 2666/20595360  
#> 3 Brazil      1999 37737/172006362  
#> 4 Brazil      2000 80488/174504898  
#> 5 China       1999 212258/1272915272  
#> 6 China       2000 213766/1280428583
```

We need to separate `rate` into two variables, `cases` and `population`

# Can we make datasets tidy?

```
separate(table3, col = rate, into = c("cases", "population"),  
          sep = "/")
```

```
#> # A tibble: 6 x 4  
#>   country    year cases population  
#>   <chr>    <int> <chr>   <chr>  
#> 1 Afghanistan 1999 745    19987071  
#> 2 Afghanistan 2000 2666   20595360  
#> 3 Brazil      1999 37737  172006362  
#> 4 Brazil      2000 80488  174504898  
#> 5 China       1999 212258 1272915272  
#> 6 China       2000 213766 1280428583
```

I've been explicit about naming all the options. R functions can work by position as well, so `separate(table3, rate, c('cases', 'population'), '/')` would work, but it's not very clear, is it?

# Can we make datasets tidy?

```
table2
```

```
#> # A tibble: 12 x 4  
#>   country    year type      count  
#>   <chr>    <int> <chr>    <int>  
#> 1 Afghanistan 1999 cases      745  
#> 2 Afghanistan 1999 population 19987071  
#> 3 Afghanistan 2000 cases      2666  
#> 4 Afghanistan 2000 population 20595360  
#> 5 Brazil      1999 cases      37737  
#> 6 Brazil      1999 population 172006362  
#> 7 Brazil      2000 cases      80488  
#> 8 Brazil      2000 population 174504898  
#> 9 China       1999 cases      212258  
#> 10 China      1999 population 1272915272  
#> 11 China      2000 cases      213766  
#> 12 China      2000 population 1280428583
```

Here there are observations on two variables in successive rows



# Can we make datasets tidy?

We need to spread these rows out into different columns

```
spread(table2, key = type, value = count)
```

```
#> # A tibble: 6 x 4
#>   country    year  cases population
#>   <chr>    <int> <int>      <int>
#> 1 Afghanistan 1999     745    19987071
#> 2 Afghanistan 2000    2666    20595360
#> 3 Brazil      1999   37737    172006362
#> 4 Brazil      2000   80488    174504898
#> 5 China       1999  212258   1272915272
#> 6 China       2000  213766   1280428583
```

country	year	key	value	country	year	cases	population
Afghanistan	1999	cases	745	Afghanistan	1999	745	19987071
Afghanistan	1999	population	19987071	Afghanistan	2000	2666	20595360
Afghanistan	2000	cases	2666	Brazil	1999	37737	172006362
Afghanistan	2000	population	20595360	Brazil	2000	80488	174504898
Brazil	1999	cases	37737	China	1999	212258	1272915272
Brazil	1999	population	172006362	China	2000	213766	1280428583
Brazil	2000	cases	80488				
Brazil	2000	population	174504898				
China	1999	cases	212258				
China	1999	population	1272915272				
China	2000	cases	213766				
China	2000	population	1280428583				

table2

# Can we make datasets tidy?

```
table4a
```

```
#> # A tibble: 3 x 3  
#>   country    `1999` `2000`  
#> * <chr>      <int> <int>  
#> 1 Afghanistan    745   2666  
#> 2 Brazil        37737  80488  
#> 3 China         212258 213766
```

Here, the variable for year is stored as a header, not as data in a cell.

We need to gather that data and put it into a column

# Can we make datasets tidy?

```
tidyr::gather(table4a, key = year, value = cases, `19
```

```
#> # A tibble: 6 x 3
#>   country    year  cases
#>   <chr>      <chr> <int>
#> 1 Afghanistan 1999     745
#> 2 Brazil      1999   37737
#> 3 China       1999  212258
#> 4 Afghanistan 2000     2666
#> 5 Brazil      2000   80488
#> 6 China       2000  213766
```

country	year	cases
Afghanistan	1999	745
Afghanistan	2000	2666
Brazil	1999	37737
Brazil	2000	80488
China	1999	212258
China	2000	213766

table4

# Making data tidy

Admittedly, spread and gather are not easy concepts, but we'll practice with them more.

1. gather collects multiple columns into 2, and only 2 columns
  - One column represents the data in the column headers
  - One column represents the values in the column
  - All other columns are repeated to keep all the data properly associated
2. spread takes two columns and makes them multiple columns
  - The values in one column form the headers to different new columns
  - The values in the other column represent the values in the corresponding cells
  - The other columns are repeated to start with, but reduce repetitions to make all associated data stay together

# Progress check

Load the data from [this link](#). You can look the structure by `head(_____)` where `_____` is what you named the dataset.

What do you think you would need to do to make this data tidy? (Hint: look at the column headers)

What function would you want to use?

Fill in the blanks:

```
gather(_____, key = _____, value = _____, _____, _____, _____,
        _____, _____, _____, _____, _____, _____)
```

This is a lot of writing. There's gotta be something simpler

05:00

# A friendly way of selecting columns

The tidyverse gives us a nice way of selecting, or not selecting columns

Instead of all the writing, we could simply say

```
pew <- read_csv('data/pew.csv')  
tidyr::gather(pew, key = income, value = count, -religion)
```

```
#> # A tibble: 180 x 3  
#>   religion      income count  
#>   <chr>      <chr>   <dbl>  
#> 1 Agnostic   <$10k     27  
#> 2 Atheist    <$10k     12  
#> 3 Buddhist   <$10k     27  
#> 4 Catholic   <$10k    418  
#> 5 Don't know/refused <$10k     15  
#> 6 Evangelical Prot <$10k    575  
#> 7 Hindu       <$10k      1  
#> 8 Historically Black Prot <$10k    228  
#> 9 Jehovah's Witness <$10k     20  
#> 10 Jewish     <$10k     19  
#> # ... with 170 more rows
```

# dp1yr verbs in the tidyverse

The dp1yr package gives us a few verbs for data manipulation

Function	Purpose
select	Select columns based on name or position
mutate	Create or change a column
filter	Extract rows based on some criteria
arrange	Re-order rows based on values of variable(s)
group_by	Split a dataset by unique values of a variable
summarize	Create summary statistics based on columns

# select

You can select columns by name or position, of course.

You can also select columns based on some criteria, which are encapsulated in functions.

- `startswith("__"), endswith("__"), contains("__")`
- `oneof("", "", "__")`

There are others; see `help(starts_with)`.



# Example

Load [this file](#). This contains daily temperature data in 2010 for some location.

```
weather <- rio::import('data/weather.csv')
# weather <- readr::read_csv(here::here('slides', 'lectures', 'data', 'FSI', 'weather.csv'))
```

```
head(weather, 2)
```

```
#>      id year month element d1 d2 d3 d4 d5 d6 d7 d8 d9 d10 d11 d12 d13
#> 1 MX17004 2010     1    tmax NA NA NA NA NA NA NA NA NA NA NA NA NA
#> 2 MX17004 2010     1    tmin NA NA NA NA NA NA NA NA NA NA NA NA NA
#>      d14 d15 d16 d17 d18 d19 d20 d21 d22 d23 d24 d25 d26 d27 d28 d29 d30 d31
#> 1  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA 27.8 NA
#> 2  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA  NA 14.5 NA
```

How would you just select the columns with the daily data?

```
select(weather, starts_with("d"))
```

# mutate

`mutate` can either transform a column in place or create a new column in a dataset

We'll use the in-built `mpg` dataset for this example, We'll select only the city and highway mileages. To use this selection later, we will need to assign it to a new name

```
mpg1 <- select(mpg, cty, hwy)
```

# mutate

We'll change the city and highway mileage to km/l from mpg. This will involve multiplying it by 1.6 and dividing by 3.8

```
mutate(mpg1, cty = cty * 1.6 / 3.8, hwy = hwy * 1.6/3.8)
```

```
#> # A tibble: 234 x 2
#>       cty    hwy
#>   <dbl> <dbl>
#> 1  7.58  12.2
#> 2  8.84  12.2
#> 3  8.42  13.1
#> 4  8.84  12.6
#> 5  6.74  10.9
#> 6  7.58  10.9
#> 7  7.58  11.4
#> 8  7.58  10.9
#> 9  6.74  10.5
#> 10 8.42  11.8
#> # ... with 224 more rows
```

This is in-place replacement

```
mutate(mpg1, cty1 = cty * 1.6/3.8, hwy1 = hwy * 1.6/3.8)
```

```
#> # A tibble: 234 x 4
#>       cty    hwy cty1 hwy1
#>   <int> <int> <dbl> <dbl>
#> 1    18    29  7.58  12.2
#> 2    21    29  8.84  12.2
#> 3    20    31  8.42  13.1
#> 4    21    30  8.84  12.6
#> 5    16    26  6.74  10.9
#> 6    18    26  7.58  10.9
#> 7    18    27  7.58  11.4
#> 8    18    26  7.58  10.9
#> 9    16    25  6.74  10.5
#> 10   20    28  8.42  11.8
#> # ... with 224 more rows
```

This creates new variables

# filter

filter extracts rows based on criteria

```
filter(mpg, cyl == 4)
```

```
#> # A tibble: 81 x 11
#>   manufacturer model displ  year  cyl trans drv   cty   hwy fl   class
#>   <chr>         <chr> <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
#> 1 audi         a4      1.8  1999    4 auto... f     18    29 p     comp...
#> 2 audi         a4      1.8  1999    4 manu... f     21    29 p     comp...
#> 3 audi         a4      2    2008    4 manu... f     20    31 p     comp...
#> 4 audi         a4      2    2008    4 auto... f     21    30 p     comp...
#> 5 audi         a4 q... 1.8  1999    4 manu... 4     18    26 p     comp...
#> 6 audi         a4 q... 1.8  1999    4 auto... 4     16    25 p     comp...
#> 7 audi         a4 q... 2    2008    4 manu... 4     20    28 p     comp...
#> 8 audi         a4 q... 2    2008    4 auto... 4     19    27 p     comp...
#> 9 chevrolet    mali... 2.4  1999    4 auto... f     19    27 r     mids...
#> 10 chevrolet   mali... 2.4  2008    4 auto... f     22    30 r     mids...
#> # ... with 71 more rows
```

This extracts only 4 cylinder vehicles

Other choices might be `cyl != 4`, `cyl > 4`, `year == 1999`, `manufacturer=="audi"`

# Exercise

We already saw the weather data. It's not tidy. Let's work to make it tidy.